


Lesson 7. Big DPs and the Curse of Dimensionality


1 Solving a Rubik's cube

- In a classic Rubik's cube, each of the 6 faces is covered by 9 stickers
- Each sticker can be one of 6 colors: white, red, blue, orange, green and yellow



- Each face of the cube can be turned independently

- Notation:



- The letter means turn the face clockwise 90°
 - ◊ For example, R means turn the right face clockwise 90°
- The letter primed means turn the face counter-clockwise 90°
 - ◊ For example, R' means turn the right face counter-clockwise 90°
- The problem: given an initial configuration of the cube, find a *shortest* sequence of turns so that each face has only one color
 - You may assume that you are allowed at most T turns
 - It turns out that any configuration can be solved in 26 turns or less: <http://cube20.org/qtm/>
- How can we formulate this problem as a dynamic program?

- Stages:


Stage $t \leftrightarrow t^{\text{th}}$ turn of the cube ($t=1, \dots, T$)

\leftrightarrow end of decision-making process ($t=T+1$)



- States in stage t (nodes): Let $\underbrace{1, \dots, N}_{\text{initial}} \underbrace{\dots, N}_{\text{solved}}$ be a list of all the possible cube configurations

Node $t_n \leftrightarrow$ being in the n^{th} configuration with turns $t, t+1, \dots, T$ remaining ($n=1, \dots, N$)

- Decisions, transitions, and rewards/costs at stage t (edges):



n^{th} configuration
 $\xrightarrow{\text{turn } x} m^{\text{th}}$ configuration
 $x \in \{R, R', U, U', L, L', D, D', F, F', B, B'\}$



- Source node:

l_1 (initial config @ turn 1)

Sink node:

end

- Shortest/longest path?

shortest

- Minimum number of turns required to solve the cube:


Length of a shortest path

- Actual sequence of turns that give the minimum number of turns to solve the cube:

Edges in the shortest path correspond to which turns to make.

2 Tetris

- You've all played Tetris before, right? Just in case...
- Tetris is a video game in which pieces fall down a 2D playing field, like this:



- Each piece is made up of four equally-sized bricks, and the playing field is 10 bricks wide and 20 bricks high
- As the pieces fall, the player can rotate them 90° in either direction, or move them left and right
- When a row is constructed without any holes, the player receives a point and the corresponding row is cleared
- The game is over once the height of bricks exceeds 20
- The problem: given a predetermined sequence of T pieces¹, determine how to place each piece in order to maximize the number of points accumulated over the course of the game
- How can we formulate this problem as a dynamic program?

¹Normally, the sequence of falling pieces is random and infinitely long. We'll consider this easier version here.

- Stages:


$\text{Stage } t \leftrightarrow \text{playing the } t^{\text{th}} \text{ piece } (t=1, \dots, T)$
 $\leftrightarrow \text{end of the decision-making process } (t=T+1)$

- States in stage t (nodes): Let $1, \dots, N$ be a list of all the possible playing fields
 $\text{empty} \leftrightarrow \text{full}$

$\text{Node } t_n \leftrightarrow \text{being in the } n^{\text{th}} \text{ playing field with pieces } t, t+1, \dots, T$
 $\text{remaining } (n=1, \dots, N)$

- Decisions, transitions, and rewards/costs at stage t (edges):

n is not a losing playing field:





n^{th} playing field

$\xrightarrow[\text{of piece } t]{\text{placement } x} m^{\text{th}}$ playing field

$x \in \text{set of all possible placements}$
 $\text{of piece } t \text{ in playing field } n$

$$y = \begin{cases} 1 & \text{if line is cleared with} \\ & \text{placement } x \text{ on playing field } n \\ & \text{w/ piece } t \\ 0 & \text{o/w} \end{cases}$$

n is a losing playing field:



- Source node:

$1, (\text{empty field @ piece 1})$

Sink node:

end

- Shortest/longest path?

longest

- Maximum number of points:

Length of a longest path

- Actual placement of pieces that give the maximum number of points:

Edges in a longest path correspond to which placements to make

3 Big DPs and the curse of dimensionality

- How big are these DPs we just formulated?
- Tetris:

- Number of states per stage: $N = 2^{200} \approx 1.61 \times 10^{60}$

- Number of stages T

⇒ Number of nodes: $N(T+1) + 1 \approx (1.61 \times 10^{60})(T+1) + 1$

- Rubik's cube:

- Number of states per stage: $N \approx 4.33 \times 10^{19}$

- Number of stages T

⇒ Number of nodes: $N(T+1) + 1 \approx (4.33 \times 10^{19})(T+1) + 1$

- The number of states is huge for both these DPs!
- ⇒ The DPs we formulated (as-is) are not solvable using today's computing power
- This is known as **the curse of dimensionality** in dynamic programming
- **Approximate dynamic programming** is an active area of research that tries to address the curse of dimensionality in various ways
 - For example, for Tetris: <https://papers.nips.cc/paper/5190-approximate-dynamic-programming-finally-performs-well-in-the-game-of-tetris.pdf>