

Lesson 15. The Birth-Death Process – Performance Measures

1 Overview

- A **birth-death process** is a Markov process with state space $\mathcal{M} = \{0, 1, 2, \dots\}$ and generator matrix

$$\mathbf{G} = \begin{bmatrix} -\lambda_0 & \lambda_0 & 0 & 0 & \cdots \\ \mu_1 & -(\lambda_1 + \mu_1) & \lambda_1 & 0 & \cdots \\ 0 & \mu_2 & -(\lambda_2 + \mu_2) & \lambda_2 & \cdots \\ \vdots & \vdots & \vdots & \vdots & \ddots \end{bmatrix}$$

- Last lesson: modeling queueing systems as birth-death processes
- This lesson:
 - How to compute steady-state probabilities
 - How to use the steady-state probabilities to compute long-run performance measures

2 Steady-state probabilities of a birth-death process

- Recall the steady-state equations:

$$\pi^\top \mathbf{G} = \mathbf{0}$$

$$\pi^\top \mathbf{1} = 1$$

- Rewrite the steady-state equations as scalar equations:

- Solve recursively in terms of π_0 :

- Define

$$d_0 = 1 \quad d_j = \frac{\lambda_0 \lambda_1 \cdots \lambda_{j-1}}{\mu_1 \mu_2 \cdots \mu_j} \quad \text{for } j = 1, 2, \dots \quad D = \sum_{i=0}^{\infty} d_i$$

- Simplify and solve for π_0 :

- Finally, solve for all the steady-state probabilities $\pi_0, \pi_1, \pi_2, \dots$:

- Note: the steady-state probabilities exist provided that $0 < \sum_{j=0}^{\infty} d_j < \infty$

Example 1 (Lesson 13 Problem 1, modified). A small ice-cream shop competes with several other ice-cream shops in a busy mall. If there are too many customers already in line at the shop, then potential customers will go elsewhere. Potential customers arrive at a rate of 20 per hour. The probability that a customer will go elsewhere is $n/3$ when there are $n \leq 3$ customers already in the system, and 1 when there are $n > 3$ customers already in the system. The server at the shop can serve customers at a rate of 10 per hour. Approximate the process of potential arrivals as Poisson, and the service times as exponentially distributed.

We can model this queueing system as a birth-death process:

$$\text{Arrival rates: } \lambda_i = \begin{cases} 20(1 - \frac{i}{3}) & \text{if } i = 0, 1, \dots, 3 \\ 0 & \text{if } i = 4, 5, \dots \end{cases} \quad \text{Service rates: } \mu_i = 10 \quad \text{if } i = 0, 1, 2, \dots$$

- Over the long run, what is the probability there are n customers in the shop? ($n = 0, 1, 2, 3$)
- Over the long run, what fraction of the time is the server busy?

3 System-level performance measures

- How do we measure the workload or congestion in the entire system?

- Expected number of customers in the system ℓ :

- Expected number of customers in the queue ℓ_q :

- s is the number of customers that can be served simultaneously
- Note that when there are $j = 0, 1, \dots, s$ customers in the system, there are 0 customers in the queue

Example 2. In Example 1, how many customers are waiting to be served (not including the customer in service)?


- Special case: λ is the arrival rate in all states, μ is the service rate of s identical servers

- **Traffic intensity** ρ :

- ◊ $\rho < 1 \Rightarrow$ the system is **stable**: customers can be served faster than they are arriving
 $\Rightarrow \rho =$ fraction of time each server is busy, a.k.a. the long-run **utilization** of each server
- ◊ $\rho \geq 1 \Rightarrow$ customers are arriving faster than they can be served

- **Offered load** o , or the expected number of busy servers:

- Relationship between ρ and ℓ_q : ℓ_q explodes as ρ approaches 1



- ◊ This example:

- constant arrival rate λ
 - constant service rate μ
 - $s = 1$
- $$\Rightarrow \ell_q = \frac{\rho^2}{1 - \rho}$$

4 Customer-level performance measures

- **Effective arrival rate** to the system λ_{eff} :

- Also referred to as the “average arrival rate”

- **Little’s law** (system-wide):

- w is the **expected waiting time**: the expected time a customer spends in the system from arrival to departure

- Intuitively, why does Little’s law hold?

- Average departure rate for an individual customer \approx

\Rightarrow Average departure rate for the whole system \approx

- System is stable \Rightarrow arrival rate = departure rate (“conservation of customers”):

- **Little’s law** (queue only):

- w_q is the **expected delay**: the expected time a customer spends in the queue

- Special case: If the service rate is a constant μ , then waiting time and delay are related like this:

Example 3. Consider the setting described in Example 1.

- What is the effective arrival rate?
- What is the expected customer delay?
- What is the expected customer waiting time?

5 Exercises

Problem 1. Consider the setting described in Example 1.

- Over the long run, at what rate are customers lost?
- Suppose that the shop makes a revenue of \$2 per customer served and pays the server \$4 per hour. What is the shop's long-run expected profit per hour?

Problem 2. Turingtown has a small urgent care center. Due to the recent increase in patients visiting the center, Turingtown is considering expanding the center. You have been asked to evaluate the current configuration.

The center has two doctors. You have estimated that the time a doctor spends with a patient is exponentially distributed with a mean of 30 minutes. Based on your observations and interviews, you have also estimated that patients arrive outside the center according to a Poisson process with a rate of 2 per hour. However, a patient will only enter the center if there is at most 1 patient already waiting (in other words, at most three patients total in the center); otherwise, the patient will opt to go to the hospital emergency room in nearby Gaussville.

- Model this setting as a birth-death process.
- Find the steady-state distribution of the number of patients in the urgent care center.
- What is the long-run expected number of patients waiting for a doctor?
- What is the long-run expected time a patient waits for a doctor?
- What fraction of the arriving patients opt to go to Gaussville instead?

Problem 3 (SMAS Exercise 8.6). The service counter at Southwest Montana Airlines has a single queue for waiting customers and two ticket agents. One of the agents is on duty at all times; the other agent goes on duty whenever the queue of customers becomes too long. Suppose that customer arrivals to the counter are well modeled as a Poisson process with rate 45 customers/hour. The agents both work at rate 30 customers/hour, and the second agent goes on duty if there are 3 or more customers at the counter (including the ones being served). Service times are exponentially distributed.

In Lesson 13, you modeled this setting as a birth-death process. Suppose the service counter area can hold at most 10 customers, including those in the queue and those being served.

- Determine the fraction of time that the second agent is on duty.
- Determine the average length of the queue, not including those being served.